In a message-passing system, what is the purpose of the "send" operation?

a) To receive messages from other processes

b) To send messages to other processes

c) To synchronize with other processes

d) To terminate a process

Which of the following is an example of a message-passing system?

a) Single-threaded application

b) Multi-threaded application

c) Distributed system

d) None of the above

What is the primary advantage of message-passing systems over shared memory systems?

a) Lower latency

b) Simplicity of programming

c) Scalability

d) Higher throughput

In a message-passing system, what is a "message queue" used for?

a) Storing data to be shared among threads

b) Storing messages to be sent to other processes

c) Synchronizing the execution of threads

d) Managing system resources

Which algorithm is commonly used for mutual exclusion in a message-passing system?

a) Peterson's algorithm

b) Semaphore algorithm

c) Lamport's algorithm

d) Amdahl's law

In Lamport's logical clock algorithm, how are events ordered?

a) By their physical timestamps

b) By their logical timestamps

c) By their sender's IP address

d) By their message size

What is the purpose of the "receive" operation in a message-passing system?

a) To acknowledge the receipt of a message

b) To retrieve messages from the message queue

c) To terminate a process

d) To broadcast messages to all processes

Which algorithm is used to ensure the total ordering of messages in a distributed system?

a) Vector clock algorithm

b) Two-phase commit protocol

c) Token ring algorithm

d) Raft consensus algorithm

Which of the following is a characteristic of a deadlock situation in a message-passing system?

a) Processes can make progress

b) Processes are blocked and cannot make progress

c) Processes are executing independently

d) Processes are waiting for input from the user

In a distributed system, what is the purpose of the "leader election" algorithm?

a) To select the process with the highest priority

b) To choose a process to coordinate the activities of others

c) To determine the order of message delivery

d) To establish a connection between processes

Answers:

b) To send messages to other processes

c) Distributed system

c) Scalability

b) Storing messages to be sent to other processes

c) Lamport's algorithm

b) By their logical timestamps

b) To retrieve messages from the message queue

a) Vector clock algorithm

b) Processes are blocked and cannot make progress

b) To choose a process to coordinate the activities of others

What is the primary goal of the Modified Flooding algorithm in a network?

a) To find the shortest path between two nodes

b) To efficiently broadcast a message to all nodes

c) To construct a minimum spanning tree

d) To establish a secure communication channel

In a Modified Flooding algorithm, how does a node decide whether to forward a message?

a) It always forwards the message.

b) It forwards the message if the sender is its immediate neighbor.

c) It forwards the message if it has not seen the message before.

d) It forwards the message only if it is the root node.

What is the purpose of constructing a DFS spanning tree with a root in a graph?

a) To find the shortest path between any two nodes.

b) To identify cycles in the graph.

c) To efficiently traverse the entire graph.

d) To establish a centralized node for broadcasting.

In a DFS spanning tree, what is the role of the root node?

a) It serves as a bridge between disconnected components.

b) It is the node with the highest degree in the graph.

c) It is the starting point for the depth-first traversal.

d) It is not connected to any other nodes in the tree.

How is the depth-first traversal of a graph different from breadth-first traversal?

a) Depth-first explores nodes level by level, while breadth-first explores branches.

b) Depth-first explores branches level by level, while breadth-first explores nodes.

c) Depth-first always starts from the root, while breadth-first can start from any node.

d) Depth-first and breadth-first are equivalent traversal methods.

In a DFS spanning tree, how are back edges used?

a) To identify tree edges

b) To identify forward edges

c) To detect cycles in the graph

d) To establish communication links

Which of the following is an advantage of using Modified Flooding in a network?

a) Reduced message overhead

b) Guaranteed delivery of messages

c) Improved network security

d) Faster convergence of routing algorithms

Answers:

b) To efficiently broadcast a message to all nodes

c) It forwards the message if it has not seen the message before.

c) To efficiently traverse the entire graph.

c) It is the starting point for the depth-first traversal.

b) Depth-first explores branches level by level, while breadth-first explores nodes.

c) To detect cycles in the graph

a) Reduced message overhead

What is the primary characteristic of a DFS spanning tree without a root in an undirected graph?

a) It always has a single central node as the root.

b) It can have multiple disjoint trees in the forest.

c) It is always a binary tree.

d) It is directed and acyclic.

In a DFS spanning tree without a root, what is the purpose of marking nodes as "visited" during the traversal?

a) To identify the root node.

b) To determine the number of nodes in the graph.

c) To prevent revisiting nodes and forming cycles.

d) To prioritize nodes for further exploration.

What term is commonly used to describe a connected component in a DFS spanning tree forest?

a) Subtree

b) Branch

c) Cluster

d) Trunk

In a DFS spanning tree without a root, how is the depth of a node defined?

a) It is the distance from the root of the tree.

b) It is the number of edges on the path from the root to the node.

c) It is the number of child nodes a node has.

d) It is always equal to 1.

How can you identify a back edge in a DFS spanning tree?

a) It connects a node to its parent in the tree.

b) It connects a node to a node in a different tree.

c) It connects a node to a node at a higher depth.

d) It connects a node to a node at a lower depth.

In a DFS spanning tree without a root, what does it mean if a node is labeled as "back edge" during the traversal?

a) It is the root node of the tree.

b) It is a leaf node with no children.

c) It is part of a cycle in the graph.

d) It is a disconnected node.

Which of the following is a common application of DFS spanning trees without a root?

a) Finding the shortest path between two nodes.

b) Detecting strongly connected components in a directed graph.

c) Constructing a minimum spanning tree.

d) Balancing load in a distributed system.

Answers:

b) It can have multiple disjoint trees in the forest.

c) To prevent revisiting nodes and forming cycles.

a) Subtree

b) It is the number of edges on the path from the root to the node.

a) It connects a node to its parent in the tree.

c) It is part of a cycle in the graph.

b) Detecting strongly connected components in a directed graph.

What is the primary goal of message ordering in a distributed system?

a) To ensure that messages are delivered in the order they were sent

b) To reduce the number of messages sent in the system

c) To improve the security of message transmission

d) To minimize network latency

Which of the following is an example of a total ordering mechanism for message delivery?

a) FIFO (First-In-First-Out)

b) Causal ordering

c) Arbitrary ordering

d) Multicast ordering

In a FIFO message ordering system, what happens when a message is delivered out of order?

a) The system reorders the message before delivery.

b) The message is discarded.

c) The message is delivered as-is.

d) The sender is notified to resend the message.

What is the main advantage of using multicast over unicast for message distribution in a network?

a) Multicast ensures that messages are delivered to all recipients in order.

b) Multicast reduces network traffic by sending a message to a group of recipients simultaneously.

c) Multicast improves message security.

d) Unicast is always faster than multicast.

Which of the following multicast delivery models ensures that messages are delivered to all members of the group in the same order?

a) Reliable multicast

b) Best-effort multicast

c) Any-order multicast

d) Causal multicast

In a distributed system, what is the purpose of a broadcast message?

a) To send a message to a specific recipient.

b) To send a message to all nodes in the network.

c) To send a message to a subset of nodes based on their roles.

d) To request acknowledgment of message delivery.

Which message delivery model is most suitable for scenarios where message loss is acceptable, and there is no strict requirement for message ordering?

a) Unicast

b) Reliable multicast

c) Best-effort multicast

d) Causal ordering

In a distributed system, what is the primary challenge in achieving global message ordering?

a) Controlling the broadcast mechanism

b) Minimizing network latency

c) Handling network failures

d) Dealing with the absence of a global clock

Answers:

a) To ensure that messages are delivered in the order they were sent

a) FIFO (First-In-First-Out)

c) The message is delivered as-is.

b) Multicast reduces network traffic by sending a message to a group of recipients simultaneously.

a) Reliable multicast

b) To send a message to all nodes in the network.

c) Best-effort multicast

d) Dealing with the absence of a global clock

In the ring-based leader election algorithm, what is the role of a node with a higher ID compared to its neighbors?

a) It becomes the leader.

b) It is removed from the ring.

c) It becomes a follower.

d) It restarts the election process.

Which statement is true for the ring-based leader election algorithm?

a) It has a time complexity of O(n).

b) It guarantees the selection of the most capable leader.

c) It requires central coordination from a master node.

d) It is resilient to network failures.

O(n log n) Algorithm:

In the O(n log n) leader election algorithm, how does it achieve a time complexity of O(n log n)?

a) By dividing the ring into smaller subrings.

b) By iterating through the ring multiple times.

c) By selecting the leader randomly.

d) By using a central coordinator.

What is an advantage of O(n log n) leader election algorithms over simple ring-based algorithms?

a) They are faster and have a lower time complexity.

b) They always elect the highest-ID node as the leader.

c) They do not require communication between nodes.

d) They are more resilient to node failures.

O(n^2) Algorithm:

Which statement is true for the O(n^2) leader election algorithm?

a) It has a time complexity of O(n log n).

b) It is a highly efficient algorithm.

c) It guarantees the selection of the lowest-ID node as the leader.

d) It is suitable for large-scale networks.

What is the primary drawback of O(n^2) leader election algorithms?

a) They are slow and impractical for large networks.

b) They are highly resilient to node failures.

c) They require central coordination.

d) They always select the highest-ID node as the leader.

Bully Algorithm:

In the Bully algorithm, when does a node initiate an election process?

a) When it has a higher ID than all other nodes.

b) When it detects that the current leader is unresponsive.

c) When it is removed from the network.

d) When a new node joins the network.

What is a potential issue with the Bully algorithm in large-scale networks?

a) It has a high time complexity.

b) It may result in many simultaneous elections.

c) It relies on central coordination.

d) It is highly efficient in terms of message complexity.

Answers:

a) It becomes the leader.

a) It has a time complexity of O(n).

a) By dividing the ring into smaller subrings.

a) They are faster and have a lower time complexity.

a) It has a time complexity of O(n log n).

a) They are slow and impractical for large networks.

b) When it detects that the current leader is unresponsive.

b) It may result in many simultaneous elections.

What is the primary goal of mutual exclusion in shared memory systems?

a) Maximizing parallelism

b) Minimizing memory usage

c) Preventing multiple processes from accessing a critical section simultaneously

d) Reducing network latency

Which of the following is a classic software-based solution for achieving mutual exclusion in shared memory systems?

a) Spinlocks

b) Semaphore

c) Mutex

d) Barrier

In the context of mutual exclusion, what is a "critical section"?

a) A section of code that executes without any restrictions

b) A section of code that must be executed by multiple processes simultaneously

c) A section of code that must be executed by only one process at a time

d) A section of code that is optional for execution

What is a disadvantage of using spinlocks for mutual exclusion?

a) High memory usage

b) High CPU utilization in busy-waiting

c) Lack of fairness in process scheduling

d) Inability to be used in multi-core systems

In a shared memory system, what is a benefit of using hardware-based solutions for mutual exclusion, such as atomic operations?

a) Lower software complexity

b) Higher memory usage

c) Improved process scheduling

d) Greater scalability in distributed systems

Which of the following problems can occur in a shared memory system if mutual exclusion is not properly enforced?

a) Deadlock

b) Starvation

c) Data race

d) Segmentation fault

What is the purpose of a mutex (mutual exclusion lock) in shared memory systems?

a) To maximize parallelism by allowing multiple processes to access critical sections concurrently

b) To coordinate communication between processes

c) To enforce mutual exclusion, ensuring only one process can access a critical section at a time

d) To perform memory management tasks

In a shared memory system, what is the primary role of a semaphore?

a) To protect critical sections from unauthorized access

b) To count the number of processes accessing shared memory

c) To provide a mechanism for distributed computing

d) To guarantee real-time scheduling for processes

Answers:

c) Preventing multiple processes from accessing a critical section simultaneously

a) Spinlocks

c) A section of code that must be executed by only one process at a time

b) High CPU utilization in busy-waiting

a) Lower software complexity

c) Data race

c) To enforce mutual exclusion, ensuring only one process can access a critical section at a time

b) To count the number of processes accessing shared memory

What is the primary goal of mutual exclusion in shared memory systems?

a) Maximizing parallelism

b) Minimizing memory usage

c) Preventing multiple processes from accessing a critical section simultaneously

d) Reducing network latency

Which of the following is a classic software-based solution for achieving mutual exclusion in shared memory systems?

a) Spinlocks

b) Semaphore

c) Mutex

d) Barrier

In the context of mutual exclusion, what is a "critical section"?

a) A section of code that executes without any restrictions

b) A section of code that must be executed by multiple processes simultaneously

c) A section of code that must be executed by only one process at a time

d) A section of code that is optional for execution

What is a disadvantage of using spinlocks for mutual exclusion?

a) High memory usage

b) High CPU utilization in busy-waiting

c) Lack of fairness in process scheduling

d) Inability to be used in multi-core systems

In a shared memory system, what is a benefit of using hardware-based solutions for mutual exclusion, such as atomic operations?

a) Lower software complexity

b) Higher memory usage

c) Improved process scheduling

d) Greater scalability in distributed systems

Which of the following problems can occur in a shared memory system if mutual exclusion is not properly enforced?

a) Deadlock

b) Starvation

c) Data race

d) Segmentation fault

What is the purpose of a mutex (mutual exclusion lock) in shared memory systems?

a) To maximize parallelism by allowing multiple processes to access critical sections concurrently

b) To coordinate communication between processes

c) To enforce mutual exclusion, ensuring only one process can access a critical section at a time

d) To perform memory management tasks

In a shared memory system, what is the primary role of a semaphore?

a) To protect critical sections from unauthorized access

b) To count the number of processes accessing shared memory

c) To provide a mechanism for distributed computing

d) To guarantee real-time scheduling for processes

Answers:

c) Preventing multiple processes from accessing a critical section simultaneously

a) Spinlocks

c) A section of code that must be executed by only one process at a time

b) High CPU utilization in busy-waiting

a) Lower software complexity

c) Data race

c) To enforce mutual exclusion, ensuring only one process can access a critical section at a time

b) To count the number of processes accessing shared memory

What is the primary purpose of token-based algorithms in distributed systems?

a) Load balancing

b) Mutual exclusion

c) Leader election

d) Message broadcasting

In a token-based algorithm, what is a "token"?

a) A small data packet that contains a message

b) A synchronization signal passed between nodes

c) A cryptographic key used for encryption

d) A data structure for storing distributed data

Which of the following problems can token-based algorithms help solve effectively?

a) Achieving perfect load balancing in a network

b) Ensuring mutual exclusion in a distributed system

c) Detecting network failures and rerouting traffic

d) Broadcasting a message to all nodes simultaneously

What is a key characteristic of a token-based algorithm with a central solution?

a) The central node has the highest priority.

b) Token passing is controlled by a central coordinator.

c) Token circulation is bidirectional.

d) All nodes have equal access to the token.

In a central token-based algorithm, what happens when a node receives the token?

a) It immediately releases the token to the next node.

b) It holds the token until it has a specific message to send.

c) It broadcasts the token to all nodes simultaneously.

d) It discards the token.

Which of the following is a potential disadvantage of central token-based algorithms?

a) High message overhead due to token circulation

b) Inefficient resource utilization

c) Complexity in managing the central coordinator

d) Scalability issues in large networks

What is the primary advantage of central token-based algorithms?

a) They are highly scalable and suitable for large networks.

b) They require minimal coordination among nodes.

c) They provide strong security through encryption.

d) They are capable of achieving perfect load balancing.

In a central token-based algorithm, what role does the central coordinator play when a node requests the token?

a) It grants the request and forwards the token.

b) It immediately releases the token to the requesting node.

c) It holds the token until a specific condition is met.

d) It selects a random node to receive the token.

Answers:

b) Mutual exclusion

b) A synchronization signal passed between nodes

b) Ensuring mutual exclusion in a distributed system

b) Token passing is controlled by a central coordinator.

a) It immediately releases the token to the next node.

c) Complexity in managing the central coordinator

d) They are capable of achieving perfect load balancing.

a) It grants the request and forwards the token.

What is the primary goal of ring-based mutual exclusion algorithms in distributed systems?

a) To maximize parallelism

b) To ensure that processes can access critical sections simultaneously

c) To prevent multiple processes from accessing a critical section simultaneously

d) To optimize network communication

In a ring-based mutual exclusion algorithm, how are processes organized in a logical ring?

a) Based on their physical location in the network

b) Based on their process ID or priority

c) In a random order

d) In a hierarchical structure

How is the token passed among processes in a ring-based mutual exclusion algorithm?

a) In a random order

b) In a predefined clockwise direction

c) In a counterclockwise direction

d) In a star topology

In a ring-based mutual exclusion algorithm, what does a process do when it has the token and wants to enter a critical section?

a) It enters the critical section immediately.

b) It sends a request message to the next process in the ring.

c) It broadcasts a message to all other processes.

d) It releases the token without entering the critical section.

What is a potential issue with ring-based mutual exclusion algorithms in the context of network failures?

a) They are highly resilient to network failures.

b) They can lead to deadlock if a process crashes.

c) They require a central coordinator to manage failures.

d) They can guarantee that no process enters the critical section.

In a ring-based mutual exclusion algorithm, what is the purpose of the "acknowledge" message?

a) It is a message that indicates a process's intention to enter the critical section.

b) It is a message sent by a process to indicate that it has successfully entered the critical section.

c) It is a message used to confirm the receipt of a request message.

d) It is a message that carries the token.

Which of the following represents a typical behavior of processes in a ring-based mutual exclusion algorithm?

a) Processes always request access to the critical section simultaneously.

b) Processes can access the critical section in any order.

c) Processes access the critical section in a predefined sequence.

d) Processes enter the critical section randomly.

In a ring-based mutual exclusion algorithm, what happens when a process receives a request from another process while it is inside the critical section?

a) It denies the request.

b) It forwards the request to the next process.

c) It grants the request and releases the token.

d) It waits until it exits the critical section to respond.

Answers:

c) To prevent multiple processes from accessing a critical section simultaneously

b) Based on their process ID or priority

b) In a predefined clockwise direction

b) It sends a request message to the next process in the ring.

b) They can lead to deadlock if a process crashes.

c) It is a message used to confirm the receipt of a request message.

c) Processes access the critical section in a predefined sequence.

d) It waits until it exits the critical section to respond.

In a non-token-based mutual exclusion algorithm, what is the primary goal?

a) Ensuring fairness in process scheduling

b) Allowing processes to access critical sections simultaneously

c) Preventing multiple processes from accessing a critical section simultaneously

d) Minimizing network latency

Which classic algorithm is an example of a non-token-based mutual exclusion algorithm?

a) Lamport's Bakery Algorithm

b) Centralized Token Algorithm

c) Ring-Based Token Algorithm

d) Bully Algorithm

In Lamport's Bakery Algorithm, how are processes assigned a number when they wish to enter a critical section?

a) Processes select a random number.

b) Processes request a number from a central coordinator.

c) Processes are assigned numbers in a sequential, incremental order.

d) Numbers are not used in this algorithm.

What is the primary purpose of the "ticket" in Lamport's Bakery Algorithm?

a) It represents a process's identity.

b) It is used as a cryptographic key for encryption.

c) It determines the order in which processes can enter the critical section.

d) It is a message exchanged between processes.

Which statement is true regarding non-token-based algorithms like Lamport's Bakery Algorithm?

a) They guarantee a fixed order in which processes enter the critical section.

b) They allow multiple processes to enter the critical section simultaneously.

c) They rely on a central coordinator to control access.

d) They have a high message overhead.

What is a drawback of non-token-based algorithms in large-scale distributed systems?

a) They are highly scalable.

b) They may lead to contention and process starvation.

c) They guarantee immediate access to the critical section.

d) They do not require coordination among processes.

In Lamport's Bakery Algorithm, what does it mean when a process's ticket number is lower than another process's ticket number?

a) The process can immediately enter the critical section.

b) The process must wait until the other process has entered the critical section.

c) The process can preempt the other process.

d) The process's request is denied.

Which factor contributes to the fairness of non-token-based algorithms like Lamport's Bakery Algorithm?

a) The order in which processes request access

b) The number of processes in the system

c) The speed of the network connection

d) The total number of critical sections in the application

Answers:

c) Preventing multiple processes from accessing a critical section simultaneously

a) Lamport's Bakery Algorithm

c) Processes are assigned numbers in a sequential, incremental order.

c) It determines the order in which processes can enter the critical section.

a) They guarantee a fixed order in which processes enter the critical section.

b) They may lead to contention and process starvation.

b) The process must wait until the other process has entered the critical section.

a) The order in which processes request access

What is the primary goal of Ricart-Agrawala's Algorithm in a distributed system?

a) To maximize parallelism by allowing processes to access critical sections simultaneously

b) To minimize network latency and message overhead

c) To ensure that only one process at a time enters a critical section

d) To achieve a fair process scheduling

In Ricart-Agrawala's Algorithm, what do processes do when they wish to enter a critical section?

a) They request permission from a central coordinator.

b) They broadcast their request to all other processes.

c) They enter the critical section immediately without requesting permission.

d) They randomly select a process to coordinate with.

How is the order in which processes request access to the critical section determined in Ricart-Agrawala's Algorithm?

a) It is determined by their process IDs.

b) It is based on a random selection process.

c) It is predefined and sequential.

d) It is based on the number of messages exchanged.

What is the purpose of the "request message" in Ricart-Agrawala's Algorithm?

a) To grant immediate access to the critical section

b) To confirm that a process has successfully entered the critical section

c) To request permission from other processes to enter the critical section

d) To serve as a token that is passed among processes

In Ricart-Agrawala's Algorithm, when can a process enter the critical section after sending request messages?

a) As soon as it sends the request messages

b) After receiving permission from all other processes

c) After receiving a specific number of acknowledgment messages

d) After a predefined waiting period

Which statement is true regarding the message overhead in Ricart-Agrawala's Algorithm?

a) It is minimal, as processes do not exchange messages.

b) It is high, as processes exchange messages with all other processes.

c) It is independent of the number of processes in the system.

d) It is constant and not affected by the number of requests.

What is a potential drawback of Ricart-Agrawala's Algorithm in the context of process failures?

a) It is highly resilient to process failures.

b) It may lead to deadlock if a process crashes.

c) It does not rely on communication among processes.

d) It guarantees that all processes can access the critical section simultaneously.

In Ricart-Agrawala's Algorithm, how do processes release access to the critical section?

a) They release access immediately upon entering the critical section.

b) They release access after receiving acknowledgment messages from other processes.

c) They release access when they are ready to leave the critical section.

d) They release access randomly.

Answers:

c) To ensure that only one process at a time enters a critical section

b) They broadcast their request to all other processes.

c) It is predefined and sequential.

c) To request permission from other processes to enter the critical section

c) After receiving a specific number of acknowledgment messages

b) It is high, as processes exchange messages with all other processes.

b) It may lead to deadlock if a process crashes.

a) They release access immediately upon entering the critical section.

What is the primary goal of Maekawa's Algorithm in a distributed system?

a) To maximize parallelism by allowing processes to access critical sections simultaneously

b) To minimize network latency and message overhead

c) To ensure that only one process at a time enters a critical section

d) To achieve a fair process scheduling

In Maekawa's Algorithm, what is the concept of "clusters"?

a) Groups of nodes in a network with the same physical location

b) Groups of processes that communicate frequently with each other

c) Logical partitions of processes that collectively grant or deny access to a critical section

d) A type of message used for communication

How are processes assigned to clusters in Maekawa's Algorithm?

a) In a completely random manner

b) Based on their process IDs

c) Based on their physical proximity in the network

d) According to a predefined mapping algorithm

What is the purpose of the "request set" in Maekawa's Algorithm?

a) To identify the critical section being accessed

b) To specify the order in which processes request access

c) To determine which clusters must be consulted for permission

d) To specify a priority level for each process

In Maekawa's Algorithm, when can a process enter the critical section after sending requests to the clusters?

a) As soon as it sends the requests

b) After receiving permission from all clusters

c) After receiving acknowledgment from one cluster

d) After a predefined waiting period

Which statement is true regarding the message overhead in Maekawa's Algorithm?

a) It is minimal, as processes do not exchange messages.

b) It is high, as processes exchange messages with all clusters.

c) It is independent of the number of processes in the system.

d) It is constant and not affected by the number of requests.

What is a potential drawback of Maekawa's Algorithm in the context of process failures?

a) It is highly resilient to process failures.

b) It may lead to deadlock if a process crashes.

c) It does not rely on communication among processes.

d) It guarantees that all processes can access the critical section simultaneously.

In Maekawa's Algorithm, how do processes release access to the critical section?

a) They release access immediately upon entering the critical section.

b) They release access after receiving acknowledgment from all clusters.

c) They release access when they are ready to leave the critical section.

d) They release access after a predefined time interval.

Answers:

c) To ensure that only one process at a time enters a critical section

c) Logical partitions of processes that collectively grant or deny access to a critical section

d) According to a predefined mapping algorithm

c) To determine which clusters must be consulted for permission

b) After receiving permission from all clusters

b) It is high, as processes exchange messages with all clusters.

b) It may lead to deadlock if a process crashes.

a) They release access immediately upon entering the critical section.